
Recurrent Neural Networks for NLP

— Intro to AI - Autumn 2025 Class —

Eleni Metheniti - eleni.metheniti@univ-tlse3.fr

Table of contents

1. [What is a neural network?](#)
2. [Recurrent Neural Networks](#)
 1. [Introduction to RNNs](#)
 2. [RNNs & NLP](#)
3. [Long Short-Term Memory NNs](#)
4. [Encoder-Decoder architecture](#)
5. [Attention Mechanism](#)
6. [Practical session: Training an RNN](#)

1. What is a neural network?

Neural Networks

- Machine learning **model**
- Learning **features** directly from data
- Types of neural networks
 - Convolutional Neural Networks (CNNs) → great for image processing
 - Recurrent Neural Networks (RNNs, LSTMs) → great for sequential data
 - Transformers → great for... everything!


Neural Networks: Neuron

- **Neuron**: the building block of a neural network
(other names: **node**, **cell**)
- Stores information, is connected to other neurons



Neural Networks: Neuron


- **Neuron**: the building block of a neural network (other names: **node**, **cell**)
- Stores information, is connected to other neurons




Yokoyama, Satoru, et al. "Cortical activation in the processing of passive sentences in L1 and L2: An fMRI study." *Neuroimage* 30.2 (2006): 570-579.



Neural Networks: Neuron



Biological neuron




Artificial neuron

Neural Networks: Neuron


Inside a neuron...

- Dot product of **input** and **weights**: the connection between two neurons
- Add the **bias**: a value that also influences the importance of the neuron
- Apply **activation function**: the *firing rate* of the neuron


Neural Networks: Activation functions



sigmoid (0 - 1)



tanh (-1 - 1)



ReLU (0 - z)

Run an activation function:

<https://developers.google.com/machine-learning/crash-course/neural-networks/activation-functions>


See more activation functions:

<https://cs231n.github.io/neural-networks-1>
<https://sebastianraschka.com/faq/docs/activation-functions.html>

Neural Networks: Layer

- **Layer:** a collection of neurons
- **The outputs of neurons become inputs to other neurons**


neurons of the same layer are **not** connected to each other...



...but each neuron is connected to **every** neuron of the next layer = fully-connected layer


Neural Network: Layer

- **Input layer:** first layer of the NN, contains **encoded data**




Neural Network: Layer

- **Hidden layer(s):** learn different aspects about the data it received




Neural Network: Layer

- **Hidden layer(s):** learn different aspects about the data it received
- One or more hidden layers, interacting with the next!



Neural Network: Layer

- **Output layer:** no function, contains values (probability of a token, label...)



e.g. binary classification
of "I hate peaches." :


0.75 → prob. of label 0 👎 = 75%

0.25 → prob. of label 1 👍 = 25%

Neural Network: Output


- **Output layer:** no function, contains values (probability of a token, label...)
- **Loss:** compare input values to predictions (= output values)

e.g. binary classification:



Neural Network: Backpropagation

- **Loss:** compare input values to predictions (= output values)
- If it's bad... **backpropagation** will tell us who is responsible! 🕵️




Neural Network: Dropout rate


Temporarily “deactivate” some nodes of the input and hidden layer(s)

probability > dropout

probability < dropout




(a) Standard Neural Net



(b) After applying dropout.


Neural Network: Epoch

- **Loss:** compare input values to predictions (= output values)
- If it's bad...
- Solution: re-adjust weights & biases, redo it all in a new **epoch**!




Neural Network: Size

- Size based on **neurons** = Count of neurons of hidden layer(s) & output layer
- Size based on **learnable parameters** = Their biases + Their weights




Structure of a NN




One fixed-size input → One fixed-size output (e.g. image classification)

Structure of a NN


one to one




one to many




many to one



many to many




many to many




One fixed-size input → Output with multiple values (e.g. image captioning)

Structure of a NN


one to one




one to many




many to one



many to many




many to many




Input with multiple values → One fixed output (e.g. text classification)

Structure of a NN


one to one




one to many




many to one



many to many





many to many



Input with multiple values → Output with multiple values (e.g. machine translation)


Neural Network: Feedforward NN

- Information flows in one direction
- (Some) drawbacks:
 - Information flows in one direction → limited contextual understanding!
 - Overfitting: good/bad with seen data, bad with unseen data
 - Requires feature engineering (see Lecture 5!)



Convolutional Neural Networks

- **Feature maps** to detect features /patterns in the input data → NN
- (Some) drawbacks for NLP:
 - Fixed input size
 - Problem with long-distance relationships



<https://domino.ai/blog/gpu-accelerated-convolutional-neural-networks-with-pytorch>

A Game of Sequential Data

- Recite the alphabet in your head...
 - In the right order!



A Game of Sequential Data

- Recite the alphabet in your head...
 - In the right order! → A, B, C, D, E, F, G, H, ...
 - Backwards!



A Game of Sequential Data

- Recite the alphabet in your head...
 - In the right order! → A, B, C, D, E, F, G, H, ...
 - Backwards! → Z, Y, X, W, V, U, ...
 - Starting from the letter “N” →




A Game of Sequential Data

- Recite the alphabet in your head...
 - In the right order! → A, B, C, D, E, F, G, H, ...
 - Backwards! → Z, Y, X, W, V, U, ...
 - Starting from the letter “N” → N, O, P, Q, R, S, T, ...
- We have learned the alphabet as a **sequence**
- Difficult in other ways!



Sequential data

- Video, audio, **text** are sequences
- But remember... 🍏🥧 ≠ 🍏📱
- We need a model that can:
 - handle context & position in the sentence
 - handle sentences with varying lengths
 - handle context left and right
 - handle long-distance relations




2. Recurrent Neural Networks

Introduction to RNNs



What is an RNN?

- Like a feedforward NN... but better with **feedback!**
- **Loops** that allow information to go back into the network




Unfolding an RNN

- **Recurrent Unit:** Neuron that holds a **hidden state** that contains the past
 - $\mathbf{x}_t$ = **input** in *time t*
 - $\mathbf{y}_t$ = **output** in *time t*
 - $\mathbf{a}_{t-1}$ = **activation** in *time t-1* = the previous activation
 - $\mathbf{a}_t$ = **activation** in *time t* = the current activation
 - $\mathbf{h}_t$ = **hidden state** of the neuron in *time t*




Unfolding an RNN

- **Backpropagation Through Time (BPTT):** “learning from its errors”
- Current output = current input + memory of past inputs




Structure of an RNN

- **Hidden state (h_t) calculation:** $h_t = \sigma(U \cdot x_t + W \cdot h_{t-1} + B)$
 - U, W = weights!
 - B = biases!



Structure of an RNN

- **Output (y_t) calculation:** $y_t = O(V \cdot h_t + C)$
 - O = activation function
 - V = weights!
 - C = biases!



RNNs: Applications

- Processing **sequential data**
 - Time-series: economics, etc.
 - NLP : words, sentences, etc. → Let's focus on this...

RNNs & NLP

NLP with RNNs: Example

Binary Classification of a sentence: "What time is it?"

-  → Category 0
-  → Category 1

NLP with RNNs: Example

- Tokenization

What time is it?


NLP with RNNs: Example

- **1st timestep:** pass the word “What” to the model

What time is it ?


NLP with RNNs: Example

- **2nd timestep:** pass the word “time” and the hidden state with “what”




NLP with RNNs: Example

- Repeat for all timesteps until end of sequence



NLP with RNNs: Example

- Binary Classification: O_5 = **Probabilities** of Category 0 🕒 & Category 1 ☀️




Vanishing gradient problem

Do you remember what was in the first image in [Slide 7](#)?

Vanishing gradient problem

Do you remember what was in the first image in [Slide 7](#)?


Human working memory has limits... and so do RNNs



Vanishing gradient problem

- **BPTT**: gradients exponentially shrink as timesteps pass
- Final timestep: early hidden states are “not important” anymore

t_2 = “what time”




t_5 = “what time is it ?”



time

Exploding gradient problem


Now you remember what was in the first image in [Slide 7](#)!



Exploding gradient problem

Now you remember what was in the first image in [Slide 7](#)!

...but it's not so important to our lecture today!



Exploding gradient problem

- What if we change the weights to make some hidden states important...?
- “time” is now very important... but this could be a problem with more data!

t_2 = “what time”

t_5 = “what time is it ?”



RNNs: Pros & Cons

- + Processing input of any length
- Computation being slow

RNNs: Pros & Cons

- + Processing input of any length
- + Modeling using **past** samples, capture dependencies!
- Computation being slow
- Cannot use any **future** input!

RNNs: Pros & Cons


- + Processing input of any length
- + Modeling using **past** samples, capture dependencies!
- + Weights are shared across time
- Computation being slow
- Cannot use any **future** input!
- Vanishing/Exploding gradients: difficulty with **long sequences**

RNNs: Pros & Cons

- + Processing input of any length
- + Modeling using **past** samples, capture dependencies!
- + Weights are shared across time
- Computation being slow
- Cannot use any **future** input!
- Vanishing/Exploding gradients: difficulty with **long sequences**



Bidirectionality

- NNs process input left → right: past → present → ~~future~~
- But what if we **reversed the sequence**?




Bidirectionality

- NNs process input left → right: past → present → ~~future~~
- But what if we **reversed the sequence**?
- The future is to the left → accessible to the model as a “past” state!




Bidirectionality: bi-RNN



<https://www.geeksforgeeks.org/bidirectional-recurrent-neural-network/>

Bidirectionality: bi-RNN



RNNs: Pros & Cons

- + Processing input of any length
- + Modeling using **past** samples, capture dependencies!
- + Weights are shared across time

- Computation being slow
- Cannot use any **future** input!

Improve the RNN architecture!


- Vanishing/Exploding gradients: difficulty with **long sequences**

3. Long Short-Term Memory Neural Networks

What is an LSTM?


- **Long Short-Term Memory** Neural Network: improvement on RNNs
- Designed to avoid vanishing/exploding gradients
- RNNs: one linear path for all previous states
- LSTMs:
 - one path for **long-term “memories”** (states) → directly influence output
 - one path for **short-term “memories”** (states) → follow the flow of information

LSTM: Cell




LSTM: Cell

Hidden State: short-term memory, **updated** by weights and biases!



LSTM: Cell

Cell State: long-term memory, **not affected** by weights and biases!




LSTM: Cell

Forget Gate:

σ = sigmoid activation function (0-1)

x = % how much long-term memory to keep




LSTM: Cell

Input Gate:

$\tanh$ = tanh activation function (-1-1)

$\times$ = % of **new memory** to update long-term memory

$+$ = add to the existing **long-term memory**




LSTM: Cell

tanh = tanh with **new long-term memory**



x = % of **new memory** to update short-term memory

Output Gate:



LSTM: Model

e.g. a bidirectional LSTM!




LSTMs: Pros & Cons

- + **Memory unit** to learn long-term dependencies in sequential data
- + **Robust** to noisy data
- + **Flexible** for many tasks
- Slower to train than RNNs because of extra calculations
- Vanishing/exploding gradients can still be a problem...
- Prone to **overfitting**

Gated Recurrent Units (GRU)

1- = inversion operation




4. Encoder-Decoder architecture

What is an encoder?

- **Part** of a neural network architecture
- Made of feedforward NNs, RNNs, LSTMs, or GRUs, etc...
- Turns a sequence into numerical representations = **context vectors**
- Pure encoder model? Yes, e.g. the sentence classifier!


What is an encoder?

1. **Encode** each input token




What is an encoder?

1. **Encode** each input token
2. 🕒 → 🕒 → 🕒: **update** hidden state based on input + context (same w&b!)



What is an encoder?

1. **Encode** each input token
2. 🕒 → 🕒 → 🕒: **update** hidden state based on input + context (same w&b!)
3. Encoder output = its final hidden state = **context vector**



Encoder's context vector

- Remember word embeddings? e.g. word2vec
- **Representation of the input sequence**
- Reminder: we use these context vectors as a plug-in for other models!
- Spoiler: we can make some very powerful word embeddings with this information about a word's context...




What is a decoder?

- **Part** of a neural network architecture
- **Reconstruct the encoded representation**
- Work **one token** at a time by picking the token with the highest probability
- **Autoregressive generation:** predicting a token based on the previously generated tokens
- Pure decoder model? Yes!




What is a decoder?

1. **Input:** the context vector as initial hidden state



What is a decoder?


1. **Input:** the context vector as initial hidden state
2. 🕒 → 🕒 → 🕒 : **predict** next token based on current hidden state + previously generated tokens




What is a decoder?

1. **Input:** the context vector as initial hidden state
2. 🕒 → 🕒 → 🕒 : **predict** next token
3. **Stop** when it predicts <eos> or max length reached!

End
of
Sequence




Encoder-decoder architecture




Je dors
↓
[Machine Translation]
↓
I am sleeping

Encoder-decoder architecture




Je dors
↓
[Machine Translation]
↓
I am sleeping

Encoder-decoder architecture




Je dors
↓
[Machine Translation]
↓
I am sleeping


Encoder-decoder architecture




Encoder-decoder architecture




Encoder-decoder architecture




Encoder-decoder architecture



Encoder-decoder architecture



Encoder-decoder architecture



Encoder-decoder: Pros & Cons

- + Better performance, adaptive
- + Generalizes across tasks
- + Adapted to NLP tasks with varying input/output lengths

Encoder-decoder: Pros & Cons


- + Better performance, adaptive
- + Generalizes across tasks
- + Adapted to NLP tasks with varying input/output lengths
- Lack of explicit alignment!
- Difficulty with long sequences!
- Over-reliance on context vector!

5. Attention mechanism



What is Attention?

- A **similarity score** between each input token (after passing encoder) and each output token (after passing through decoder)
- **Attention score's** goal: highest scores between related tokens, e.g. $Je \leftrightarrow I$
- **Attention mechanism**: a NN in-between Encoder-Decoder!


Adding Attention to seq2seq




Adding Attention to seq2seq




Adding Attention to seq2seq




Adding Attention to seq2seq




Inside the context vector...




Adding Attention to seq2seq




Inside the context vector...



Adding Attention to seq2seq




Inside the context vector...




The (ideal) attention weights






But is it enough attention?

- Vaswani et al. (2017): **Self-attention** = foundation for transformers & LLMs!
- Goodbye seq2seq with attention...



But is it enough attention?

- Vaswani et al. (2017): **Self-attention** = foundation for transformers & LLMs!
- Hello **transformer with self-attention**!



Thank you for your attention (weights)!

Let's head to Google Colab!

<https://colab.research.google.com/drive/1MWqyE46KXhKfL-hpQBMhAHYSIsKejtal?usp=sharing>



Practical session refresher

- **Index:** the number assigned to a token, e.g. 'ajout' $\leftrightarrow$ 23
- **Padding:** adding 0s to the sentence vector after <eos>, in order to make all our input sequences the same length (0 = ignored)
- **Dropout:** deactivate temporarily some nodes
- **Epoch:** one full training of the NN with all the data
 - **Forward/backward (pass):** one processing pass of the data in one direction
- **Batch:** how many sentences we will process simultaneously
- **Loss:** function that compares the target and predicted output values, aka how good/bad the predictions are!

Further resources

- [A Neural Network Playground](#): visualize how neural networks look inside!
- [Google tutorials on Neural networks](#) ([en français](#))
-  (Reliable) reading material:
 - [The Unreasonable Effectiveness of Recurrent Neural Networks](#)
 - [Visualizing A Neural Machine Translation Model](#) - [en français](#)
 - [Dive into Deep Learning](#) - everything NN: book, code, math (advanced!)
-  YouTube tutorials:
 - [3Blue1Brown](#)
 - [StatQuest](#)
 - [The AI hacker](#)